

Data Science Workshop-1

ITER, SOA University

Centre for Data Science
SOA University



Contents

- Handling Missing Data, filtering out, Filling in missing data
- Data transformation, Removing duplicates, Transforming using a function, Replacing values, renaming index
- Discretization and binning
- Detecting and filtering outliers
- Permutation and random sampling
- Extension data types
- String Manipulation
- Categorical Data



Data Preaparation

- Data preparation: loading, cleaning, transforming, and rearranging



Handling Missing Data

- The goals of pandas is to make working with missing data as painless as possible
- The isna method gives us a Boolean Series with True where values are null
- The built-in Python None value(key-word) is also treated as NA
- notna Negation of isna, returns True for non-NA values and False for NA values.



isna and notna

```
import pandas as pd
import numpy as np
data=pd.Series([1.2, -3.5, np.nan, 0, None])
data
```

```
0    1.2
1   -3.5
2    NaN
3    0.0
4    NaN
dtype: float64
```

data.isna()		data.notna()	
0	False	0	True
1	False	1	True
2	True	2	False
3	False	3	True
4	True	4	False
dtype: bool		dtype: bool	



dropna

- To drop nan values, use `notna()` and boolean indexing

```
data[data.notna()]
```

```
0    1.2  
1   -3.5  
3    0.0  
dtype: float64
```

- `dropna()`: it returns the Series with only the nonnull data and

```
data.dropna()
```

```
0    1.2  
1   -3.5  
3    0.0  
dtype: float64
```

corresponding index values.



dropna(axis)

- If columns which contain missing values are to be removed, use axis=1 or axis='columns'
- Default value of axis is 0 or 'index'

```
import pandas as pd
import numpy as np
data=pd.DataFrame([[1.2,-3.5,np.nan,4],
                   [np.nan,0,7,None],
                   [4.7,7.4,-0.7,0.4]])
```

data

Out[4]:

	0	1	2	3
0	1.2	-3.5	NaN	4.0
1	NaN	0.0	7.0	NaN
2	4.7	7.4	-0.7	0.4

```
data.dropna(axis=1)
```

Out[5]:

	1
0	-3.5
1	0.0
2	7.4



dropna(how)

- row or column to be removed from DataFrame, when we have at least one NaN or all NaN

```
data.dropna(how='all')
```

Out[7]:

	0	1	2	3
0	1.2	-3.5	NaN	4.0
1	NaN	0.0	7.0	NaN
2	4.7	7.4	-0.7	0.4

```
data.dropna(how='any')
```

Out[8]:

	0	1	2	3
2	4.7	7.4	-0.7	0.4



dropna(thresh)

- keep only rows containing at least a certain number of non-null values
- thresh=2, means at least two Non-null values should be there in the row.

data

	0	1	2	3
0	1.2	-3.5	NaN	4.0
1	NaN	0.7	0.0	NaN
2	4.7	7.4	-0.7	0.4

data.dropna(thresh=4)

	0	1	2	3
2	4.7	7.4	-0.7	0.4

data.dropna(thresh=3)

	0	1	2	3
0	1.2	-3.5	NaN	4.0
2	4.7	7.4	-0.7	0.4

data.dropna(thresh=2)

	0	1	2	3
0	1.2	-3.5	NaN	4.0
1	NaN	0.7	0.0	NaN
2	4.7	7.4	-0.7	0.4



Filling Nan values

- Fill NA/NaN values using the specified method.
- Calling fillna with a constant replaces missing values with that value
- We can specify particular values for particular columns, by passing a dict of values to the value keyword.

```
data.fillna(99)
```

```
out[29]:
```

	A	B	C	D
0	1.2	-3.5	99.0	4.0
1	99.0	0.0	7.0	99.0
2	4.7	7.4	-0.7	0.4

```
dic={'A':99,'B':999,'C':9999}  
data.fillna(value=dic)
```

```
out[31]:
```

	A	B	C	D
0	1.2	-3.5	9999.0	4.0
1	99.0	0.0	7.0	NaN
2	4.7	7.4	-0.7	0.4



fillna(method,limit)

- Method to use for filling NaN values
- ffill: propagate last valid observation forward to next valid.
- bfill: use next valid observation to fill gap.
- limit: If method is specified, this is the maximum number of consecutive NaN values to forward/backward fill.
- If method is not specified, this is the maximum number of entries along a column where NaNs will be filled.

```
data.fillna(method='ffill')
```

Out[37]:

	A	B	C	D
0	1.2	-3.5	NaN	4.0
1	1.2	0.0	7.0	4.0
2	4.7	7.4	-0.7	0.4

```
data.fillna(0,limit=2)
```

Out[69]:

	A	B	C	D	E
0	1.2	4.7	-3.5	0.0	4.0
1	0.0	0.0	0.0	4.7	7.4
2	0.0	0.0	0.0	7.0	0.0
3	NaN	4.7	7.4	-0.7	0.4

```
data.fillna(method='ffill',limit=2)
```

Out[70]:

	A	B	C	D	E
0	1.2	4.7	-3.5	NaN	4.0
1	1.2	4.7	0.0	4.7	7.4
2	1.2	4.7	0.0	7.0	7.4
3	NaN	4.7	7.4	-0.7	0.4



Data Transformation

- Duplicate rows may be found in a DataFrame
- The method `data.duplicated()` returns a Boolean Series indicating whether each row is a duplicate (its column values already appeared) or not
- `drop_duplicates()` returns a DataFrame with rows where the duplicated array is False

```
data = pd.DataFrame(  
    {"c1": ["a", "b", "a", "b", "a", "b"],  
     "c2": [1, 1, 1, 2, 2, 2]})  
data
```

out[75]:

	c1	c2
0	a	1
1	b	1
2	a	1
3	b	2
4	a	2
5	b	2

```
data.duplicated()
```

out[76]:

```
0    False  
1    False  
2     True  
3    False  
4    False  
5     True  
dtype: bool
```

```
data.drop_duplicates()
```

out[77]:

	c1	c2
0	a	1
1	b	1
3	b	2
4	a	2



drop_duplicates(subset,keep)

- subset: Only consider certain columns for identifying duplicates, by default use all of the columns.
- Keep: Determines which duplicates to keep.
 - 'first' : Drop duplicates except for the first occurrence.
 - 'last' : Drop duplicates except for the last occurrence.

```
data['c3']=['!','@','#','$','#','@']  
data.drop_duplicates('c3')
```

```
data.drop_duplicates(  
    subset=['c3'],keep='first')
```

```
data.drop_duplicates(  
    subset=['c3'],keep='last')
```

	c1	c2	c3
0	a	1	!
1	b	1	@
2	a	1	#
3	b	2	\$

	c1	c2	c3
0	a	1	!
1	b	1	@
2	a	1	#
3	b	2	\$

	c1	c2	c3
0	a	1	!
3	b	2	\$
4	a	2	#
5	b	2	@



apply, map, applymap

- apply works on a row / column of a DataFrame or a series

```
import pandas as pd
import numpy as np
frame = pd.DataFrame(np.random.randn(4, 3), columns=list('abc'), index=['w', 'x', 'y', 'z'])
frame
```

	a	b	c
w	0.187124	-0.978202	-0.915644
x	0.621087	-0.181380	-0.474292
y	-0.222500	-0.050965	-0.386923
z	-0.303024	0.365252	1.381760

```
col = lambda x: x.max()
frame.apply(col)
```

```
a    0.621087
b    0.365252
c    1.381760
dtype: float64
```



apply, map, applymap

- map works element-wise on a Series(A dataframe column can be thought of as a series)
- applymap works element-wise on a DataFrame(The reason for the name applymap is to differentiate from map function for series)

```
ele = lambda x: x**2  
frame.applymap(ele)
```

```
frame['a'].map(ele)
```

	a	b	c
w	0.035015	0.956880	0.838405
x	0.385749	0.032899	0.224953
y	0.049506	0.002597	0.149710
z	0.091824	0.133409	1.909261

```
W    0.035015  
X    0.385749  
Y    0.049506  
Z    0.091824
```

```
Name: a, dtype: float64
```



Check Point

- Create the following data frame
- Find the person with maximum height
- Create a BMI column($BMI = Kg/(m^2)$)
- Add a category column, mentioning obessed if $BMI > 20$, else Normal

	Name	Height	Weight
0	Salman	1.68	78
1	Aiswarya	1.63	55
2	Shahid	1.71	72
3	Kareena	1.65	53



Solution to above que

```
frame[frame['Height']==max(frame['Height'])]['Name']
```

```
2    Shahid
```



Solution to above que

```
frame[frame['Height']==max(frame['Height'])]['Name']
```

```
2    Shahid
```

```
ele = lambda x: x**2  
x=frame['Height'].map(ele)  
y=frame['Weight']  
frame['BMI']=y/x  
frame
```



Solution to above que

```
frame[frame['Height']==max(frame['Height'])]['Name']
```

2 Shahid

```
ele = lambda x: x**2
x=frame['Height'].map(ele)
y=frame['Weight']
frame['BMI']=y/x
frame
OW=lambda x: 'Obessed' if x>20 else 'Normal'
frame['Category']=frame['BMI'].map(OW)
frame
```

	Name	Height	Weight	BMI	Category
0	Salman	1.68	78	27.636054	Obessed
1	Aiswarya	1.63	55	20.700817	Obessed
2	Shahid	1.71	72	24.622961	Obessed
3	Kareena	1.65	53	19.467401	Normal



map,apply,applymap

- map method accepts a dictionary object containing a mapping to do the transformation of values

```
Name_Gen={'Salman':'Male','Aiswarya':'Female',  
          'Shahid':'Male','Kareena':'Female'}  
frame["Gender"] = frame["Name"].map(Name_Gen)  
frame
```

	Name	Height	Weight	BMI	Category	Gender
0	Salman	1.68	78	27.636054	Obsessed	Male
1	Aiswarya	1.63	55	20.700817	Obsessed	Female
2	Shahid	1.71	72	24.622961	Obsessed	Male
3	Kareena	1.65	53	19.467401	Normal	Female



replace

- `df.replace(to_replace,value)`: Replace values given in `to_replace` with `value`.
- `to_replace`: can be a number or string, or list of values, can also be a dictionary to specify different replacement values for different existing values.
- `value`: can be a scalar, list matching the length with `to_replace`, can also be a dict of values can be used to specify which value to use for each column



replace

	c1	c2	c3
0	a	a	a
1	a	1	!
2	b	1	@
3	a	1	#
4	b	2	\$
5	a	2	#
6	b	2	@

```
data.replace(to_replace="a"  
            ,value='*')
```

	c1	c2	c3
0	*	*	*
1	*	1	!
2	b	1	@
3	*	1	#
4	b	2	\$
5	*	2	#
6	b	2	@

```
data.replace(to_replace  
            =[ "a", 'b' ],  
            value='*')
```

	c1	c2	c3
0	*	*	*
1	*	1	!
2	*	1	@
3	*	1	#
4	*	2	\$
5	*	2	#
6	*	2	@



Frame Title

```
data.replace({"a": '*',  
             'b': '&'})
```

```
data.replace(to_replace  
            =["a", 'b'],  
            value=['*', '&'])
```

```
data.replace(to_replace=  
            {'c1': {'a': '*',  
                    'b': '&'}})
```

	c1	c2	c3
0	*	*	*
1	*	1	!
2	&	1	@
3	*	1	#
4	&	2	\$
5	*	2	#
6	&	2	@

	c1	c2	c3
0	*	*	*
1	*	1	!
2	&	1	@
3	*	1	#
4	&	2	\$
5	*	2	#
6	&	2	@

	c1	c2	c3
0	*	a	a
1	*	1	!
2	&	1	@
3	*	1	#
4	&	2	\$
5	*	2	#
6	&	2	@



Renaming axis labels

- Axis labels can be transformed by a function or mapping to produce new, differently labeled objects.

```
import pandas as pd
import numpy as np
data=pd.DataFrame(
    {'c1':['a','a','b','a','b','a','b'],
     'c2':['a',1,1,1,2,2,2],
     'c3':['a','!', '@', '#', '$', '#', '@']},
    index=['d','s','f','g','h','j','k'])
data
```

Out[17]:

	c1	c2	c3
d	a	a	a
s	a	1	!
f	b	1	@
g	a	1	#
h	b	2	\$
j	a	2	#
k	b	2	@

```
def transform(x):
    return x.upper()
data.index.map(transform)
```

Out[19]:

Index(['D', 'S', 'F', 'G', 'H', 'J', 'K'], dtype='object')

In [18]:

```
data.index = data.index.map(transform)
data
```

Out[18]:

	c1	c2	c3
D	a	a	a
S	a	1	!
F	b	1	@
G	a	1	#
H	b	2	\$
J	a	2	#
K	b	2	@



Renaming axis labels

- Use rename to create a transformed version of a dataset without modifying the original
- DataFrame.rename supports two calling conventions
 - (index=index_mapper, columns=columns_mapper)
 - (mapper, axis='index', 'columns')

```
data.rename(index=str.title,  
            columns=str.upper)
```

Out[16]:

	C1	C2	C3
D	a	a	a
S	a	1	!
F	b	1	@
G	a	1	#
H	b	2	\$
J	a	2	#
K	b	2	@

```
data.rename(str.lower, axis='columns')
```

Out[20]:

	c1	c2	c3
D	a	a	a
S	a	1	!
F	b	1	@
G	a	1	#
H	b	2	\$
J	a	2	#
K	b	2	@

```
data.rename(str.lower, axis='index')
```

Out[21]:

	c1	c2	c3
d	a	a	a
s	a	1	!
f	b	1	@
g	a	1	#
h	b	2	\$
j	a	2	#
k	b	2	@

Discretization

- `pd.cut` : when we need to segment and sort data values into bins.
- useful for going from a continuous variable to a categorical variable.
For example, `cut` could convert ages to groups of age ranges
- `pandas.cut(x, bins, right, labels, precision)`
 - `x`: array to be binned. Must be 1-dimensional
 - `bins`: The criteria to bin by. It can be an int or sequence of scalars, or `IntervalIndex`
 - `int` : Defines the number of equal-width bins in the range of `x`
 - `sequence of scalars` : Defines the bin edges allowing for non-uniform width
 - `IntervalIndex` : Defines the exact bins to be used
 - `right`: Indicates whether bins includes the rightmost edge or not. If `right == True` (the default), then the bins `[1, 2, 3, 4]` indicate `(1,2]`, `(2,3]`, `(3,4]` .
 - `labels`: Specifies the labels for the returned bins.
 - `precision`: The precision at which to store and display the bins labels.
`precision=2` option limits the decimal precision to two digits



Discretization

```
x=np.array([1,2,3,7,4])  
pd.cut(x,bins=2)
```

```
[(0.994, 4.0], (0.994, 4.0], (0.994, 4.0], (4.0, 7.0], (0.994, 4.0]]  
Categories (2, interval[float64, right]): [(0.994, 4.0] < (4.0, 7.0]]
```

```
x=np.array([1,2,3,7,4])  
y=[0,2,6]  
pd.cut(x,bins=y)
```

```
[(0.0, 2.0], (0.0, 2.0], (2.0, 6.0], NaN, (2.0, 6.0]]  
Categories (2, interval[int64, right]): [(0, 2] < (2, 6]]
```

```
bins = pd.IntervalIndex.from_tuples([(0, 2),(2,3),(3,7)])  
pd.cut(x, bins)
```

```
[(0, 2], (0, 2], (2, 3], (3, 7], (3, 7]]  
Categories (3, interval[int64, right]): [(0, 2] < (2, 3] < (3, 7]]
```



Discretization

```
x=np.array([1,2,3,7,4])  
y=[0,2,6]  
pd.cut(x,bins=y,right=False)
```

```
[[0.0, 2.0), [2.0, 6.0), [2.0, 6.0), NaN, [2.0, 6.0)]  
Categories (2, interval[int64, left]): [[0, 2) < [2, 6)]
```

```
data = np.random.uniform(size=20)  
pd.cut(data, 4, precision=2)
```

```
[(0.66, 0.88], (0.66, 0.88], (0.0068, 0.23], (0.45, 0.66], (0.66, 0.88], ...,  
(0.45, 0.66], (0.0068, 0.23], (0.45, 0.66], (0.66, 0.88], (0.66, 0.88)]  
Length: 20  
Categories (4, interval[float64, right]): [(0.0068, 0.23] < (0.23, 0.45] < (0.4  
5, 0.66] < (0.66, 0.88]]
```

```
ls=["bad", "medium", "good", 'Best']  
pd.cut(data, 4, precision=2, labels=ls)
```

```
['Best', 'Best', 'bad', 'good', 'Best', ..., 'good', 'bad', 'good', 'Best', 'Be  
st']  
Length: 20  
Categories (4, object): ['bad' < 'medium' < 'good' < 'Best']
```



Discretizations

- The object pandas returns is a special Categorical object.
- This has two attributes codes, and categories

```
data = np.random.uniform(size=20)
ls=["bad", "medium", "good", 'Best']
x=pd.cut(data, 4, precision=2, labels=ls)
x.codes
```

```
array([1, 2, 3, 0, 1, 3, 0, 1, 2, 2, 3, 1, 0, 2, 1, 3, 2, 1, 2, 1],
      dtype=int8)
```

```
x.categories
```

```
Index(['bad', 'medium', 'good', 'Best'], dtype='object')
```

- `pd.value_counts()`: The bin counts for the result of `pandas.cut`

```
pd.value_counts(x)
```

```
medium    7
good      6
Best      4
bad       3
dtype: int64
```



pd.qcut()

- Quantile-based discretization function.
- `pd.qcut(x,q)`:
 - `x`: 1d array or series
 - `q`: int or list-like of float Number of quantiles. 10 for deciles, 4 for quartiles. Alternatively array of quantiles, e.g. `[0, .25, .5, .75, 1.]` for quartiles.

```
data = np.random.standard_normal(1000)
quartiles = pd.qcut(data, 4, precision=2)
quartiles

[(-0.56, 3.09], (-2.69, -0.75], (0.56, 3.09], (-0.063, 0.56], (-0.75, -0.063],
..., (0.56, 3.09], (-0.75, -0.063], (-0.063, 0.56], (-0.063, 0.56], (-0.75, -0.
063]]
Length: 1000
Categories (4, interval[float64, right]): [(-2.69, -0.75] < (-0.75, -0.063] <
(-0.063, 0.56] < (0.56, 3.09]]
```

```
pd.value_counts(quartiles)
```

```
(-2.69, -0.75]      250
(-0.75, -0.063]     250
(-0.063, 0.56]      250
(0.56, 3.09]        250
dtype: int64
```



Detecting outliers

- Create a 1000 by 4 matrix, entries from standard normal distribution



Detecting outliers

- Create a 1000 by 4 matrix, entries from standard normal distribution(`np.random.standard_normal((1000, 4))`)
- find the minimum, maximum, mean, standard deviation, 25percentile, 50percentile and 75percentile.



Detecting outliers

- Create a 1000 by 4 matrix, entries from standard normal distribution(`np.random.standard_normal((1000, 4))`)
- find the minimum, maximum, mean, standard deviation, 25percentile, 50percentile and 75percentile.(`df.describe()`)
- select all rows having a value exceeding 3 or -3



Detecting outliers

- Create a 1000 by 4 matrix, entries from standard normal distribution(`np.random.standard_normal((1000, 4))`)
- find the minimum, maximum, mean, standard deviation, 25percentile, 50percentile and 75percentile.`(df.describe())`
- select all rows having a value exceeding 3 or -3 `data[(data.abs())>3].any(axis="columns")`
- `np.sign(data)` produces 1 and -1 values based on whether the values in data are positive or negative



Permutation and random sampling

- Permuting (randomly reordering) a numpy array can be done using `numpy.random.permutation` function.
- one way for permuting dataframes is you permute the row indices or column indices then you pass that to `df.iloc` or `df.take()`

```
df = pd.DataFrame(np.arange(5 * 7).reshape((5, 7)))  
df
```

	0	1	2	3	4	5	6
0	0	1	2	3	4	5	6
1	7	8	9	10	11	12	13
2	14	15	16	17	18	19	20
3	21	22	23	24	25	26	27
4	28	29	30	31	32	33	34

```
sampler = np.random.permutation(5)  
sampler
```



Permutation and random sampling

```
df.take(sampler)
```

	0	1	2	3	4	5	6
1	7	8	9	10	11	12	13
0	0	1	2	3	4	5	6
3	21	22	23	24	25	26	27
4	28	29	30	31	32	33	34
2	14	15	16	17	18	19	20

```
df.iloc[sampler]
```

	0	1	2	3	4	5	6
1	7	8	9	10	11	12	13
0	0	1	2	3	4	5	6
3	21	22	23	24	25	26	27
4	28	29	30	31	32	33	34
2	14	15	16	17	18	19	20



Permutation and random sampling

- By invoking take with axis="columns", we could also select a permutation of the columns
- To select a random subset without replacement (the same row cannot appear twice), you can use the sample method
- a sample with replacement (to allow repeat choices), pass replace=True to sample

```
column_sampler = np.random.permutation(7)
column_sampler
df.take(column_sampler, axis="columns")
```

```
df.sample(n=3)
```

```
df.sample(n=3, replace=True)
```

	0	5	3	1	4	6	2
0	0	5	3	1	4	6	2
1	7	12	10	8	11	13	9
2	14	19	17	15	18	20	16
3	21	26	24	22	25	27	23
4	28	33	31	29	32	34	30

	0	1	2	3	4	5	6
2	14	15	16	17	18	19	20
0	0	1	2	3	4	5	6
1	7	8	9	10	11	12	13

	0	1	2	3	4	5	6
4	28	29	30	31	32	33	34
3	21	22	23	24	25	26	27
4	28	29	30	31	32	33	34

Computing Indicator/Dummy Variables

- converting a categorical variable into a dummy or indicator matrix
- If a column in a DataFrame has k distinct values, you would derive a matrix with k columns containing all 1s and 0s
- pandas has a `pandas.get_dummies` function for doing this
- `pandas.get_dummies(data, prefix, dtype)`
 - data: Data of which to get dummy indicators
 - prefix: String to append DataFrame column names
 - dtype: Data type for new columns



Dummy Variable

```
import pandas as pd
import numpy as np
df=pd.DataFrame({'C1':['A','B','C','A','C'],
                 'C2':range(5)})
df
```

	C1	C2
0	A	0
1	B	1
2	C	2
3	A	3
4	C	4

```
pd.get_dummies(df['C1']) pd.get_dummies(df['C1'], dtype=float)
```

	A	B	C
0	1	0	0
1	0	1	0
2	0	0	1
3	1	0	0
4	0	0	1

	A	B	C
0	1.0	0.0	0.0
1	0.0	1.0	0.0
2	0.0	0.0	1.0
3	1.0	0.0	0.0
4	0.0	0.0	1.0

```
pd.get_dummies(df["C1"], dtype=float,prefix='C1')
```

```
dummy=pd.get_dummies(df["C1"], dtype=float,prefix='C1')
df_with_dummy = df[["C2"]].join(dummy)
df_with_dummy
```

	C1_A	C1_B	C1_C
0	1.0	0.0	0.0
1	0.0	1.0	0.0

	C2	C1_A	C1_B	C1_C
0	0	1.0	0.0	0.0
1	1	0.0	1.0	0.0

Check point

- Collect 10 samples from the uniform distribution in the interval $[0,1]$
- create 4 bins $[0,0.25)$, $[0.25,0.5)$, $[0.5,0.75)$, $[0.75,1)$, and keep all elements into the respective beans
- create 4 dummy columns, with names as above intervals. If the first element belongs to a particular interval, in the respective column it should be True, and False otherwise

```
values = np.random.uniform(size=10)
bins = [0, 0.25, 0.5, 0.75, 1]
pd.get_dummies(pd.cut(values, bins))
```

- Change the code, to get the exact output



Extension types

- pandas was originally built upon the capabilities present in NumPy
- Many pandas concepts, such as missing data, were implemented using what was available in NumPy
- Building on NumPy led to a number of shortcomings
 - when missing data was introduced into such data, pandas converted the data type to float64 and used np.nan to represent null values
 - Datasets with a lot of string data were computationally expensive
 - Some data types, like time intervals, timedeltas, and timestamps with time zones, could not be supported efficiently without using computationally expensive arrays of Python objects.
- More recently, pandas has developed an extension type system allowing for new data types to be added even if they are not supported natively by NumPy.



Extension types

```
: s = pd.Series([1, 2, 3, None])
s
0    1.0
1    2.0
2    3.0
3    NaN
dtype: float64
```

```
: s.dtype
dtype('float64')
```

```
: s = pd.Series([1, 2, 3, None], dtype=pd.Int64Dtype())
s
0      1
1      2
2      3
3    <NA>
dtype: Int64
```

```
: s.dtype
Int64Dtype()
```



Extension types

Pandas extension types provide a way to extend the functionality of pandas by creating custom data types. This can offer several advantages:

- **Efficient Storage:** Extension types can be more memory-efficient, allowing you to represent data in a more compact form.
- **Improved Code Readability:** Creating custom extension types can lead to more readable and self-explanatory code
- **Domain-Specific Data Handling:** Extension types enable you to create data types tailored to specific domains
- **Compatibility with Pandas Ecosystem:** Extension types can be seamlessly integrated into the broader pandas ecosystem, ensuring compatibility with various libraries and tools that work with pandas data structures.



String Manipulation

- Built-In String Object Methods(Split, strip,join, count, find, replace etc, see Strings PPT)
- Regular Expressions(Regular expressions provide a flexible way to search or match string patterns in text see next few slides)
- String Functions in pandas(For a string operation on pandas objects)
 - String and regular expression methods can be applied (passing a lambda or other function) to each value using data.map, but it will fail on the NA (null) values.
 - To cope with this, pandas Series objects has has str attribute, which skip over NA values.



regular expression

- RegEx can be used to check if a string contains the specified search pattern.
- The re module offers a set of functions that allows us to search a string for a match:
 - search Returns a Match object if there is a match anywhere in the string
 - there is more than one match, only the first occurrence of the match will be returned
 - A Match Object is an object containing information about the search and the result.
 - The Match object has properties and methods used to retrieve information about the search, and the result:
 - .span() returns a tuple containing the start-, and end positions of the match.
 - .string returns the string passed into the function
 - .group() returns the part of the string where there was a match
 - findall Returns a list containing all matches



regular expression

```
match=re.search('Hi','Helli Himanshu')  
match
```

```
<re.Match object; span=(6, 8), match='Hi'>
```

```
match.group()
```

```
'Hi'
```

Character range
from A to Z

Character range
from a to z

Indicates a digit
from 0 to 9

[A-Za-z]{2}\d{3}

Square brackets
containing
character
range

It means exactly 2
occurrences of any
character from
preceding pattern

It means exactly 3
occurrences of any
character from
preceding pattern

e.g., CS229, cs231

Examples that match above pattern



regular expression

```
s='SOACS219ITEREE223'  
import re  
r='[A-Za-z]{2}\d{3}'  
match=re.search(r,s)  
match.group()
```

'CS219'

```
re.findall(r,s)
```

['CS219', 'EE223']

- Can you guess the output of the following code?

```
my_str = '''Hi my name is Ashis and gmail address is  
ashisinmath@gmail.com and soa email is ashispati@soa.ac.in'''  
r="([\w_.-]+)@([\w]+\.[\w]+)"  
#r="([a-zA-Z0-9_.-]+)@([a-zA-Z0-9-]+\.[a-zA-Z0-9-]+)"  
re.findall(r, my_str)
```



string functions in pandas

- Vectorized string functions for Series and Index.
- **NAs** stay **NA** unless handled otherwise by a particular method.

```
data = {"Ashis": "ashisinmath@gmail.com",  
        "Modi": "iamnarendra@gmail.com",  
        "Manu": "manoj@soa.ic.in", "Naveen": np.nan}  
data = pd.Series(data)  
data
```

```
Ashis      ashisinmath@gmail.com  
Modi       iamnarendra@gmail.com  
Manu       manoj@soa.ic.in  
Naveen     NaN  
dtype: object
```

```
data.str.contains("gmail")
```

```
Ashis      True  
Modi       True  
Manu       False  
Naveen     NaN  
dtype: object
```

```
data.str[:5]
```

```
Ashis      ashis  
Modi       iamna  
Manu       manoj  
Naveen     NaN  
dtype: object
```

```
data.str.split("@")
```

```
Ashis      [ashisinmath, gmail.com]  
Modi       [iamnarendra, gmail.com]  
Manu       [manoj, soa.ic.in]  
Naveen     NaN  
dtype: object
```



Categorical data

- A categorical variable takes on a fixed, number of possible values (categories). Example gender variable.
- Categoricals are a pandas data type corresponding to categorical variables in statistics
- the data structure consists of a categories array and an integer array of codes which point to the real value in the categories array.

